

Introducing Mechanistic Interpretability: Demistify black boxes with **Circuit Analysis**¹ & **Monosemanticity**²

J. Setpal

February 1, 2024



**MACHINE LEARNING
@ PURDUE**

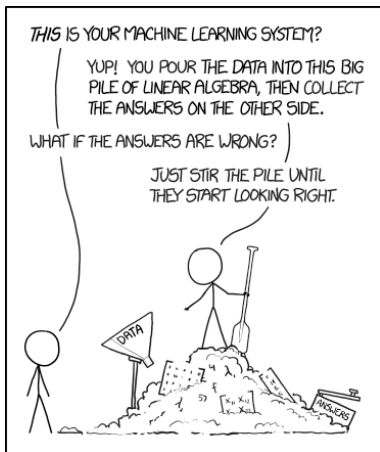
¹<https://transformer-circuits.pub/2021/framework/>

²<https://transformer-circuits.pub/2023/monosemantic-features/>

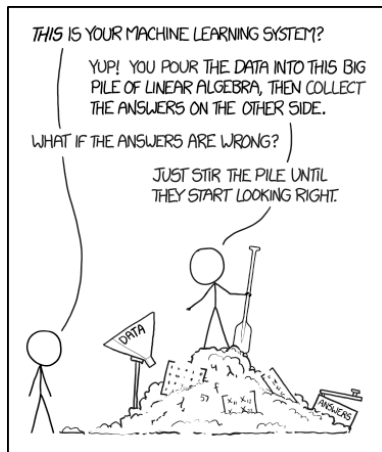
- ① Background & Intuition
- ② Transformer Circuit Analysis
- ③ Towards Monosemanticity

- ① Background & Intuition
- ② Transformer Circuit Analysis
- ③ Towards Monosemanticity

What is Interpretability?

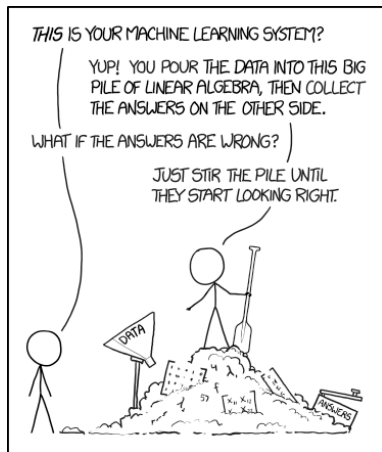


What is Interpretability?



Interpretability within Machine Learning is the **degree** to which we can understand the **cause** of a decision, and use it to consistently predict the model's prediction.

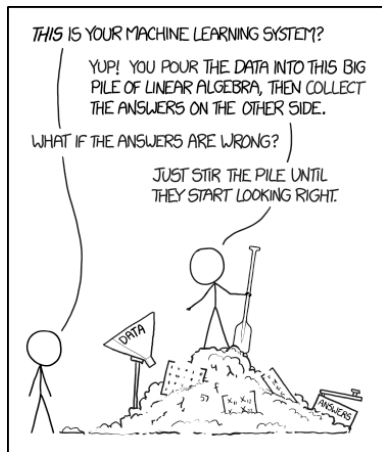
What is Interpretability?



Interpretability within Machine Learning is the **degree** to which we can understand the **cause** of a decision, and use it to consistently predict the model's prediction.

This is easy for shallow learning.

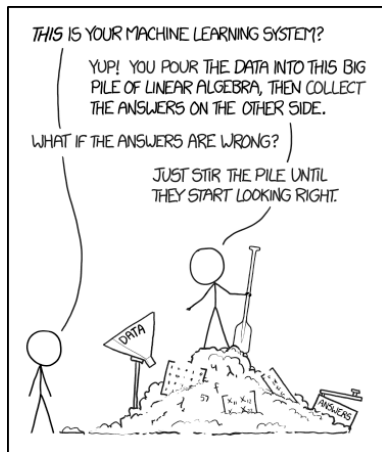
What is Interpretability?



Interpretability within Machine Learning is the **degree** to which we can understand the **cause** of a decision, and use it to consistently predict the model's prediction.

This is easy for shallow learning. For deep learning however, it is a **lot harder.**

What is Interpretability?




Interpretability within Machine Learning is the **degree** to which we can understand the **cause** of a decision, and use it to consistently predict the model's prediction.

This is easy for shallow learning. For deep learning however, it is a **lot harder.**

Today, we will interpret deep neural networks (transformer).


What will we Achieve Today?



Specifically, we'll analyze the 1-layer attention model.

For mathematical simplicity, this model ignores biases, layer-norm and dense layers.

What will we Achieve Today?




Specifically, we'll analyze the 1-layer attention model.

For mathematical simplicity, this model ignores biases, layer-norm and dense layers.

Why is this useful?

What will we Achieve Today?



Specifically, we'll analyze the 1-layer attention model.

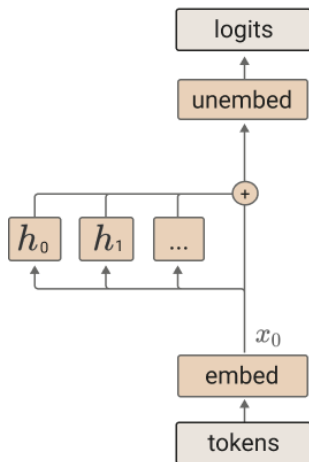
For mathematical simplicity, this model ignores biases, layer-norm and dense layers.

Why is this useful?

If we are able to *completely understand* a toy model, we can:

- understand why attention works.

What will we Achieve Today?



Specifically, we'll analyze the 1-layer attention model.

For mathematical simplicity, this model ignores biases, layer-norm and dense layers.

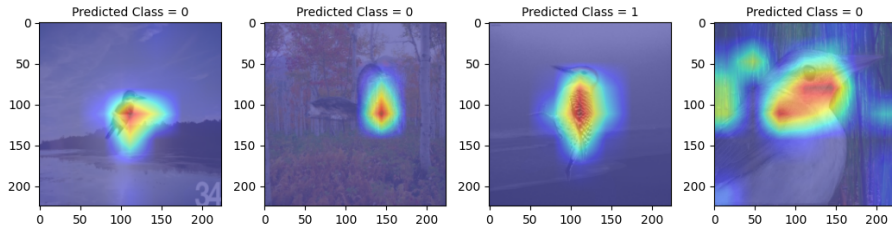
Why is this useful?

If we are able to *completely understand* a toy model, we can:

- understand why attention works.
- observe recurring patterns in complex models.


What is *Mechanistic* Interpretability?

Most of interpretability seeks to extract representations from weights:



What is *Mechanistic* Interpretability?


Most of interpretability seeks to extract representations from weights:



Mechanistic Interpretability is a subset of interpretability, that places a focus on **reverse engineering neural networks**.

What is *Mechanistic* Interpretability?

Most of interpretability seeks to extract representations from weights:

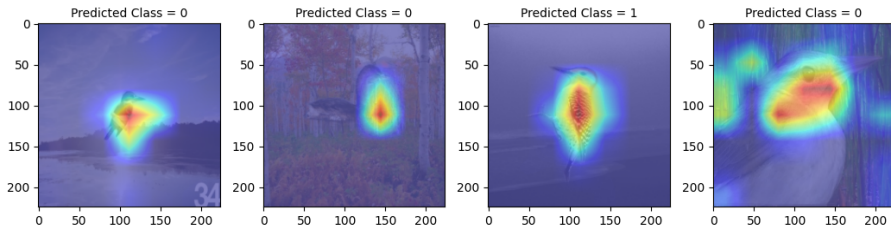


Mechanistic Interpretability is a subset of interpretability, that places a focus on **reverse engineering neural networks**.

It seeks to understand functions that *individual neurons* play in the inference of a neural network.

What is *Mechanistic* Interpretability?

Most of interpretability seeks to extract representations from weights:



Mechanistic Interpretability is a subset of interpretability, that places a focus on **reverse engineering neural networks**.

It seeks to understand functions that *individual neurons* play in the inference of a neural network.

This can subsequently be used to offer high-level explanations for decisions, as well as guarantees during inference.

Outline

- ① Background & Intuition
- ② Transformer Circuit Analysis
- ③ Towards Monosemanticity

Self-Attention Synopsis

n -gram models used the following incorrect assumption:

$$p(x_t | \{x_i\}_{i=1}^{t-1}; \theta) \not\approx p(x_t | x_{t-1}; \theta) \quad (1)$$

Self-Attention Synopsis

n -gram models used the following incorrect assumption:

$$p(x_t | \{x_i\}_{i=1}^{t-1}; \theta) \not\approx p(x_t | x_{t-1}; \theta) \quad (1)$$

Why $\not\approx$?

Self-Attention Synopsis

n -gram models used the following incorrect assumption:

$$p(x_t | \{x_i\}_{i=1}^{t-1}; \theta) \not\approx p(x_t | x_{t-1}; \theta) \quad (1)$$

Why $\not\approx$? It's because context is important!

Self-Attention Synopsis

n -gram models used the following incorrect assumption:

$$p(x_t | \{x_i\}_{i=1}^{t-1}; \theta) \not\approx p(x_t | x_{t-1}; \theta) \quad (1)$$

Why $\not\approx$? It's because context is important!

But, so is *efficiency*. Self-Attention solves this by effectively creating a **trainable database**.

Self-Attention Synopsis

n -gram models used the following incorrect assumption:

$$p(x_t | \{x_i\}_{i=1}^{t-1}; \theta) \not\approx p(x_t | x_{t-1}; \theta) \quad (1)$$

Why $\not\approx$? It's because context is important!

But, so is *efficiency*. Self-Attention solves this by effectively creating a **trainable database**.

We query it to subset the important tokens. For $\{x_i\}_{i=1}^t$,

$$\alpha_i = \sigma_{softmax} \left(\frac{q_i k_i^T}{\sqrt{d_k}} \right) \quad (2)$$

$$(3)$$

Where q_i, k_i, v_i are each independent parameter matrices.

Self-Attention Synopsis

n -gram models used the following incorrect assumption:

$$p(x_t | \{x_i\}_{i=1}^{t-1}; \theta) \not\approx p(x_t | x_{t-1}; \theta) \quad (1)$$

Why $\not\approx$? It's because context is important!

But, so is *efficiency*. Self-Attention solves this by effectively creating a **trainable database**.

We query it to subset the important tokens. For $\{x_i\}_{i=1}^t$,

$$\alpha_i = \sigma_{softmax} \left(\frac{q_i k_i^T}{\sqrt{d_k}} \right) \quad (2)$$

$$h(x) = \sum_{i=1}^t \alpha_i v_i \quad (3)$$

Where q_i, k_i, v_i are each independent parameter matrices.

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

(5)

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

$$= (A \otimes W_O W_V) \cdot x \quad (5)$$

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

$$= (A \otimes W_O W_V) \cdot x \quad (5)$$

The *disjointed* nature of A , $W_O W_V$ tells us a lot!

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

$$= (A \otimes W_O W_V) \cdot x \quad (5)$$

The *disjointed* nature of A , $W_O W_V$ tells us a lot!

- a. A and $W_O W_V$ are fundamentally independent entities.

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

$$= (A \otimes W_O W_V) \cdot x \quad (5)$$

The *disjointed* nature of A , $W_O W_V$ tells us a lot!

- a. A and $W_O W_V$ are fundamentally independent entities.
- b. A describes which token information moves through, $W_O W_V$ describes which residual subspace to read from and write to.

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

$$= (A \otimes W_O W_V) \cdot x \quad (5)$$

The *disjointed* nature of A , $W_O W_V$ tells us a lot!

- A and $W_O W_V$ are fundamentally independent entities.
- A describes which token information moves through, $W_O W_V$ describes which residual subspace to read from and write to.

$$MHA(x_0) = x_0 + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot x_0 \quad (6)$$

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

$$= (A \otimes W_O W_V) \cdot x \quad (5)$$

The *disjointed* nature of A , $W_O W_V$ tells us a lot!

- A and $W_O W_V$ are fundamentally independent entities.
- A describes which token information moves through, $W_O W_V$ describes which residual subspace to read from and write to.

$$MHA(x_0) = x_0 + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot x_0 \quad (6)$$

Our final transformer has the following equation:

$$T(t_0) = (I \otimes W_U) \cdot MHA((I \otimes W_E) \cdot t_0) \quad (7)$$

Reframing using Tensorization (1/3)

We can represent attention using **tensor products**:

$$h(x) = (I \otimes W_O) \cdot (A \otimes I) \cdot (I \otimes W_V) \cdot x \quad (4)$$

$$= (A \otimes W_O W_V) \cdot x \quad (5)$$

The *disjointed* nature of A , $W_O W_V$ tells us a lot!

- A and $W_O W_V$ are fundamentally independent entities.
- A describes which token information moves through, $W_O W_V$ describes which residual subspace to read from and write to.

$$MHA(x_0) = x_0 + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot x_0 \quad (6)$$

Our final transformer has the following equation:

$$T(t_0) = (I \otimes W_U) \cdot MHA((I \otimes W_E) \cdot t_0) \quad (7)$$

Why is this important?

Reframing using Tensorization (2/3)

We begin by simplifying to just T :

$$T = (I \otimes W_U) \cdot MHA(I \otimes W_E) \tag{8}$$

(9)

(10)

Reframing using Tensorization (2/3)

We begin by simplifying to just T :

$$T = (I \otimes W_U) \cdot MHA(I \otimes W_E) \quad (8)$$

$$= (I \otimes W_U) \cdot (I \otimes W_E + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot I \otimes W_E) \quad (9)$$

$$(10)$$

Reframing using Tensorization (2/3)

We begin by simplifying to just T :

$$T = (I \otimes W_U) \cdot MHA(I \otimes W_E) \quad (8)$$

$$= (I \otimes W_U) \cdot (I \otimes W_E + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot I \otimes W_E) \quad (9)$$

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

Reframing using Tensorization (2/3)

We begin by simplifying to just T :

$$T = (I \otimes W_U) \cdot MHA(I \otimes W_E) \quad (8)$$

$$= (I \otimes W_U) \cdot (I \otimes W_E + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot I \otimes W_E) \quad (9)$$

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

Here's the breakdown:

- a. $W_U W_E$ approximate bigram statistics.

Reframing using Tensorization (2/3)

We begin by simplifying to just T :

$$T = (I \otimes W_U) \cdot MHA(I \otimes W_E) \quad (8)$$

$$= (I \otimes W_U) \cdot (I \otimes W_E + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot I \otimes W_E) \quad (9)$$

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

Here's the breakdown:

- $W_U W_E$ approximate bigram statistics.
- A^h dictates where the attention heads attend.

Reframing using Tensorization (2/3)

We begin by simplifying to just T :

$$T = (I \otimes W_U) \cdot MHA(I \otimes W_E) \quad (8)$$

$$= (I \otimes W_U) \cdot (I \otimes W_E + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot I \otimes W_E) \quad (9)$$

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

Here's the breakdown:

- a. $W_U W_E$ approximate bigram statistics.
- b. A^h dictates where the attention heads attend.
- c. $W_U W_O^h W_V^h W_E$ describes the **behavior of logits if we attend to a given token**.

Reframing using Tensorization (2/3)

We begin by simplifying to just T :

$$T = (I \otimes W_U) \cdot MHA(I \otimes W_E) \quad (8)$$

$$= (I \otimes W_U) \cdot (I \otimes W_E + \sum_{h \in H} (A^h \otimes W_O^h W_V^h) \cdot I \otimes W_E) \quad (9)$$

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

Here's the breakdown:

- $W_U W_E$ approximate bigram statistics.
- A^h dictates where the attention heads attend.
- $W_U W_O^h W_V^h W_E$ describes the **behavior of logits if we attend to a given token**.

Observation: The equation is linear, if we fix attention patterns.

Reframing using Tensorization (3/3)

Finally, let's also unpack attention in **tensor-product form**.

Reframing using Tensorization (3/3)

Finally, let's also unpack attention in **tensor-product form**.

First, we can display key-value matrix operations:

$$q_i = (I \otimes W_Q W_E) \cdot t_0 \quad (11)$$

$$k_i = (I \otimes W_K W_E) \cdot t_0 \quad (12)$$

Reframing using Tensorization (3/3)

Finally, let's also unpack attention in **tensor-product form**.

First, we can display key-value matrix operations:

$$q_i = (I \otimes W_Q W_E) \cdot t_0 \quad (11)$$

$$k_i = (I \otimes W_K W_E) \cdot t_0 \quad (12)$$

And then apply them to unnormalized³ attention:

$$A = \sigma_{softmax} \left([q_i k_j^T]_{i,j} \right) \quad (13)$$

$$= \sigma_{softmax} \left(t_0^T \cdot (I \otimes W_E^T W_Q^T) \cdot (I \otimes W_K W_E) \cdot t_0 \right) \quad (14)$$

$$= \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

³to ease computation.

Unravelling QK, OV Circuits (1/3)

Here's the two tensor equations combined:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

Unravelling QK, OV Circuits (1/3)

Here's the two tensor equations combined:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

Q: Is there anything interesting about these two? (similarities, differences)

Unravelling QK, OV Circuits (1/3)

Here's the two tensor equations combined:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

Q: Is there anything interesting about these two? (similarities, differences)

Here's my observations:

- a. It's a much *simpler* recomposition of feedforward inference.

Unravelling QK, OV Circuits (1/3)

Here's the two tensor equations combined:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

Q: Is there anything interesting about these two? (similarities, differences)

Here's my observations:

- It's a much *simpler* recomposition of feedforward inference.
- A is the *only* non-linear operation.

Unravelling QK, OV Circuits (1/3)

Here's the two tensor equations combined:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

Q: Is there anything interesting about these two? (similarities, differences)

Here's my observations:

- It's a much *simpler* recomposition of feedforward inference.
- A is the *only* non-linear operation.
- A **learns independently** from the rest of the tensor equation.

Unravelling QK, OV Circuits (1/3)

Here's the two tensor equations combined:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

Q: Is there anything interesting about these two? (similarities, differences)

Here's my observations:

- It's a much *simpler* recomposition of feedforward inference.
- A is the *only* non-linear operation.
- A **learns independently** from the rest of the tensor equation.

However, we're still missing one.

Unravelling QK, OV Circuits (2/3)

Importantly, both equations have $(|voc|, |voc|)$ size matrices:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

These chained tensor operations are our **circuits**, and lie at the heart of the transformer architecture.

Unravelling QK, OV Circuits (2/3)

Importantly, both equations have $(|voc|, |voc|)$ size matrices:

$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

These chained tensor operations are our **circuits**, and lie at the heart of the transformer architecture.

- a. The **Output-Value(OV) Circuit** $W_U W_O^h W_V^h W_E$: determines how attending to a token affects logits.

Unravelling QK, OV Circuits (2/3)

Importantly, both equations have $(|voc|, |voc|)$ size matrices:

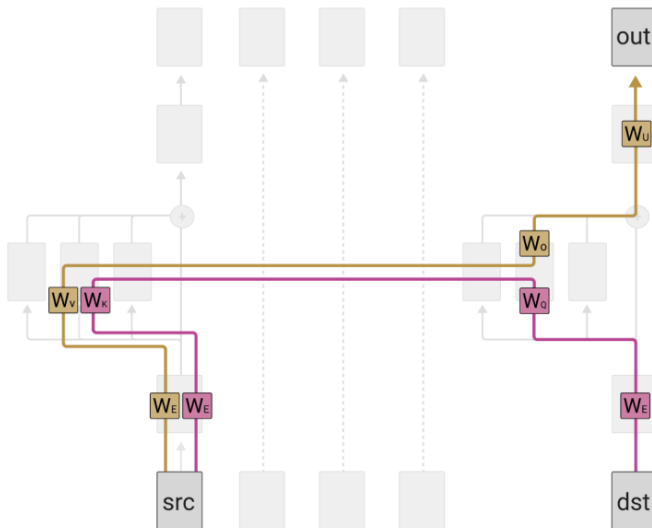
$$T = W_U W_E + \sum_{h \in H} (A^h \otimes W_U W_O^h W_V^h W_E) \quad (10)$$

$$A = \sigma_{softmax} \left(t_0^T \cdot W_E^T W_Q^T W_K W_E \cdot t_0 \right) \quad (15)$$

These chained tensor operations are our **circuits**, and lie at the heart of the transformer architecture.

- The **Output-Value(OV) Circuit** $W_U W_O^h W_V^h W_E$: determines how attending to a token affects logits.
- The **Query-Key(QK) Circuit** $W_E^T W_Q^T W_K W_E$: determines which tokens to attend to.

Unravelling QK, OV Circuits (3/3)



Interpretation as Skip-Trigrams

We can think through inference procedure with *single* source token.⁴

⁴for simplicity.

Interpretation as Skip-Trigrams

We can think through inference procedure with *single* source token.⁴

From there, we look at the largest QK and OV entries.

Some examples of large entries QK/OV circuit

Source Token	Destination Token	Out Token	Example Skip Tri-grams
" perfect"	" are", " looks", " is", " provides"	" perfect", " super", " absolute", " pure"	" perfect... are perfect", " perfect... looks super"
" large"	" contains", " using", " specify", " contain"	" large", " small", " very", " huge"	" large... using large", " large... contains small"
" two"	" One", "\n ", " has", "\r\n ", "One"	" two", " three", " four", " five", " one"	" two... One two", " two... has three"
"lambda"	" \$\\", " }\\", " +\\", " (\\", " \${\""	" lambda", "sorted", " lambda", "operator"	" lambda... \$\lambda", " lambda... +\lambda"
"nbsp"	" &", " \&", " }&", ">&", " =&"	" nbsp", "01", "gt", "00012", " nbs", "quot"	" nbsp... nbsp", " nbsp... > nbsp"
"Great"	"The", " The", " the", " contains", " /"	" Great", " great", " poor", " Every"	"Great... The Great", "Great... the great"

⁴for simplicity.

Eigenvalue Analysis

Most of the prominent behaviours include copying. We can identify this using **eigenvalue analysis**.

Eigenvalue Analysis

Most of the prominent behaviours include copying. We can identify this using **eigenvalue analysis**. Recall from the definition of eigenvectors,

$$Wv = \lambda v; \lambda \in \mathbb{C} \quad (16)$$

This is useful when we map a vector space upon itself.

Eigenvalue Analysis

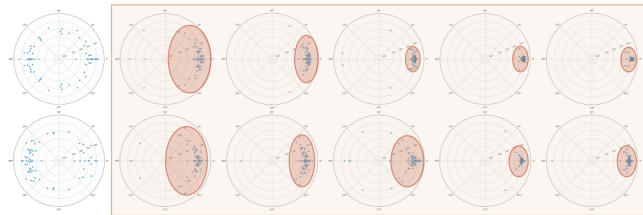
Most of the prominent behaviours include copying. We can identify this using **eigenvalue analysis**. Recall from the definition of eigenvectors,

$$Wv = \lambda v; \lambda \in \mathbb{C} \quad (16)$$

This is useful when we map a vector space upon itself.

Eigenvalue analysis of **first layer** attention head OV circuits

10/12 of layer 1 heads have mostly positive OV eigenvalues, and appear to significantly perform copying



← non-positive eigenvalues
not copying heads?

positive eigenvalues
copying heads? →

We use a **log scale** to represent magnitude, since it varies by many orders of magnitude.

Eigenvalue distribution for randomly initialized weights. Note that the mostly – and in some cases, entirely – positive eigenvalues we observe are very different from what we randomly expect.



Importantly, note that positive eigenvalues mean they are copying 'on average', and are not definitive.

- ① Background & Intuition
- ② Transformer Circuit Analysis
- ③ Towards Monosemanticity**

Problem Setup

Q: Is anyone familiar with the the curse of dimensionality?

Problem Setup

Q: Is anyone familiar with the the curse of dimensionality?

A: For NNs, basically latent space $\propto |\text{layers}|^c$.

This makes them tough to analyze at scale.

Problem Setup

Q: Is anyone familiar with the the curse of dimensionality?

A: For NNs, basically latent space $\propto |\text{layers}|^c$.

This makes them tough to analyze at scale. In addition, models are *incredibly efficient* at information compression.

Problem Setup


Q: Is anyone familiar with the the curse of dimensionality?

A: For NNs, basically latent space $\propto |\text{layers}|^c$.


This makes them tough to analyze at scale. In addition, models are *incredibly efficient* at information compression.

This is **superposition**.

HYPOTHETICAL DISENTANGLED MODEL



OBSERVED MODEL



Problem Setup

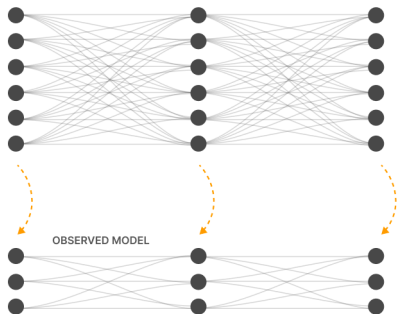
Q: Is anyone familiar with the the curse of dimensionality?

A: For NNs, basically latent space $\propto |\text{layers}|^c$.

This makes them tough to analyze at scale. In addition, models are *incredibly efficient* at information compression.

This is **superposition**.

HYPOTHETICAL DISENTANGLED MODEL



When we perform an individual analysis of neurons, it fires for unrelated concepts.

This is **polysematicity**.

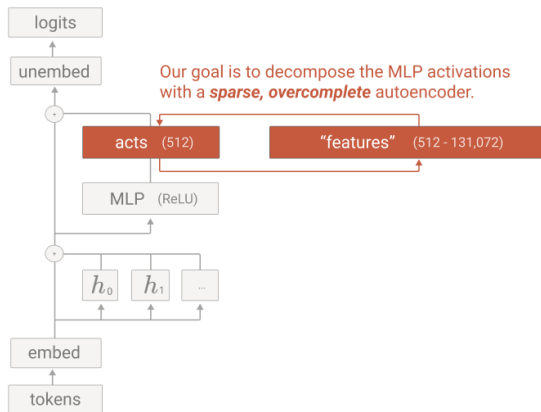
Updated Architecture

Previously, we used an **attention-only** model, since the MLP was too hard to analyze mathematically.

Updated Architecture

Previously, we used an **attention-only** model, since the MLP was too hard to analyze mathematically.

Let's instead analyze the following architecture *empirically*:



Training Setup

	Transformer	Sparse Autoencoder
Layers	1 Attention Block 1 MLP Block	1 ReLU 1 Linear
MLP Size	512	$512 \times f \in \{1, \dots, 256\}^5$
Dataset	The Pile (100B tokens)	Activations (8B samples)
Loss	Autoregressive Log-Likelihood	$L2$ Reconstruction $L1$ on hidden-layer activation

⁵ $f = 8$ for our analysis

Training Setup

	Transformer	Sparse Autoencoder
Layers	1 Attention Block 1 MLP Block	1 ReLU 1 Linear
MLP Size	512	$512 \times f \in \{1, \dots, 256\}^5$
Dataset	The Pile (100B tokens)	Activations (8B samples)
Loss	Autoregressive Log-Likelihood	$L2$ Reconstruction $L1$ on hidden-layer activation

Objective: *polysemantic activations* \xrightarrow{Tr} **monosemantic features**.

⁵ $f = 8$ for our analysis

Training Setup

	Transformer	Sparse Autoencoder
Layers	1 Attention Block 1 MLP Block	1 ReLU 1 Linear
MLP Size	512	$512 \times f \in \{1, \dots, 256\}^5$
Dataset	The Pile (100B tokens)	Activations (8B samples)
Loss	Autoregressive Log-Likelihood	L2 Reconstruction L1 on hidden-layer activation

Objective: *polysemantic activations* \xrightarrow{Tr} **monosemantic features**.

The sparse, overcomplete autoencoder is trained against this objective.

1. **Sparse** because we constraint activations (L1 penalty).
2. **Overcomplete** because the hidden layer exceeds the input dimension.

⁵ $f = 8$ for our analysis

Sparse Dictionary Learning

We can motivate our objective transformation by linear factorization:

$$x^j \approx b + \sum_i f_i(x^j) d_i \quad (17)$$

$$f_i = \sigma_{ReLU}(W_E(x - b_D) + b_E) \quad (18)$$

where d_i is the 'feature direction' represented as columns of the W_D .

Sparse Dictionary Learning

We can motivate our objective transformation by linear factorization:

$$x^j \approx b + \sum_i f_i(x^j) d_i \quad (17)$$

$$f_i = \sigma_{ReLU}(W_E(x - b_D) + b_E) \quad (18)$$

where d_i is the 'feature direction' represented as columns of the W_D .

Some interesting implementation notes:

- a. Training data \propto interpretable features.

Sparse Dictionary Learning

We can motivate our objective transformation by linear factorization:

$$x^j \approx b + \sum_i f_i(x^j) d_i \quad (17)$$

$$f_i = \sigma_{ReLU}(W_E(x - b_D) + b_E) \quad (18)$$

where d_i is the 'feature direction' represented as columns of the W_D .

Some interesting implementation notes:

- Training data \propto interpretable features.
- Tying b_D before the encoder and after the decoder improves performance.

Sparse Dictionary Learning

We can motivate our objective transformation by linear factorization:

$$x^j \approx b + \sum_i f_i(x^j) d_i \quad (17)$$

$$f_i = \sigma_{\text{ReLU}}(W_E(x - b_D) + b_E) \quad (18)$$

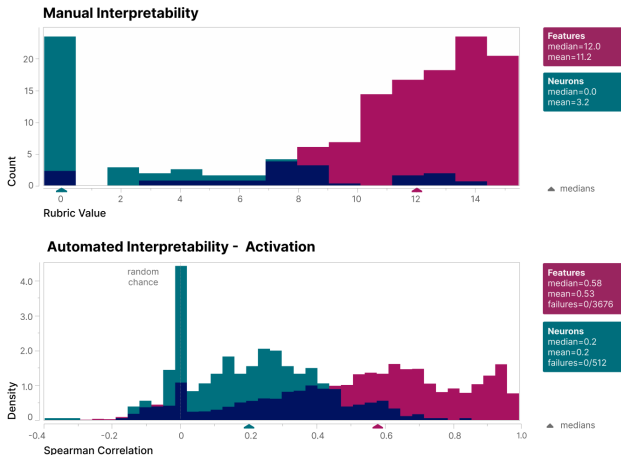
where d_i is the 'feature direction' represented as columns of the W_D .

Some interesting implementation notes:

- Training data \propto interpretable features.
- Tying b_D before the encoder and after the decoder improves performance.
- Dead neurons are periodically *resampled* to improve feature representations.

Evaluating Interpretability


Reliable evaluations on interpretability were scored based on a rubric:



Features were found to be interpretable when score > 8 .


Conditional Sampling

They can be used to steer generation.



Conditional Sampling


They can be used to steer generation.



Approach: Set high values of features demonstrating desired behaviors, and then sample from the model.

Conditional Sampling

They can be used to steer generation.



Approach: Set high values of features demonstrating desired behaviors, and then sample from the model.

We observe that interpreted features are actively used by the model.

If you can view this screen, I am making a mistake.

Thank you!

Have an awesome rest of your day!

Slides: <https://cs.purdue.edu/homes/jsetpal/slides/mechinterp.pdf>